

Lecture 14: Shortest Paths

Tim LaRock

larock.t@northeastern.edu

bit.ly/cs3000syllabus

Business


Last time: Betweenness Centrality

Betweenness centrality is used as a proxy for the importance of a node in facilitating connections between other nodes.

For node u , betweenness is measured as the ratio of shortest paths between all other pairs of nodes (s, t) that u lies on. Formally:

$$B(u) = \sum_{s \neq t \neq u} \frac{\sigma_{st}(u)}{\sigma_{st}}$$

Where σ_{st} is the number of shortest paths between nodes s and t and $\sigma_{st}(u)$ is the number of those shortest paths that include u .



Last time: How do we compute shortest paths?

To compute betweenness centrality, we need to compute shortest paths for all pairs of nodes.

Rather than jumping straight there, let's first solve a simpler problem: finding the length of the shortest path from a single node s to all other nodes in the graph (called the *single source shortest path* problem)

We can use BFS!

Last time: Single Source Shortest Paths with BFS

$\text{dist}[v]$ stores the current estimate of the distance between our source node s and the node v , initialized to infinity

We walk along every edge of the graph and check whether the distance currently stored for v could be made shorter by routing through u

If yes, we update the distance, and store u as the *predecessor* (similar to parent) of v in the shortest path

Once BFS is done, we have the shortest path length from s to every node v stored in $\text{dist}[v]$ and we can recover a shortest path for any node by following pred back to s

SSSP-BFS(s):

$\text{dist}[u] \leftarrow \infty$ for all $u \in V$

$\text{pred}[u] \leftarrow \text{null}$ for all $u \in V$

$\text{dist}[s] \leftarrow 0$

$Q \leftarrow s$

While Q is not empty:

$u \leftarrow \text{Pull}(Q)$


For $v \in \text{Neighbors}(u)$:

If $\text{dist}[v] > \text{dist}[u] + 1$:

$\text{dist}[v] = \text{dist}[u] + 1$

$\text{pred}[v] = u$

$\text{Push}(Q, v)$



Last time: Single Source Shortest Paths with BFS

$\text{dist}[v]$ stores the current estimate of the distance between our source node s and the node v , initialized to infinity

We walk along every edge of the graph and check whether the distance currently stored for v could be made shorter by routing through u

If yes, we update the distance, and store u as the *predecessor* (similar to parent) of v in the shortest path

Once BFS is done, we have the shortest path length from s to every node v stored in $\text{dist}[v]$ and **we can recover a shortest path for any node by following pred back to s**

SSSP-BFS(s):

$\text{dist}[u] \leftarrow \infty$ for all $u \in V$

$\text{pred}[u] \leftarrow \text{null}$ for all $u \in V$

$\text{dist}[s] \leftarrow 0$

$Q \leftarrow s$

While Q is not empty:

$u \leftarrow \text{Pull}(Q)$

For $v \in \text{Neighbors}(u)$:


If $\text{dist}[v] > \text{dist}[u] + 1$:

$\text{dist}[v] = \text{dist}[u] + 1$

$\text{pred}[v] = u$

Push(Q, v)

When we compute betweenness centrality, we will need all of the paths! How?



Recovering all paths from SSSP-BFS

A simple modification allows us to store all of the possible shortest paths.


We just need to adjust our $\text{pred}[u]$ data structure to store a *list* of predecessors, rather than just one!

```
SSSP-BFS( $s$ ):  
   $\text{dist}[u] \leftarrow \infty$  for all  $u \in V$   
   $\text{pred}[u] \leftarrow \text{null}$  for all  $u \in V$   
   $\text{dist}[s] \leftarrow 0$   
   $Q \leftarrow s$   
  While  $Q$  is not empty:  
     $u \leftarrow \text{Pull}(Q)$   
    For  $v \in \text{Neighbors}(u)$ :  
      If  $\text{dist}[v] > \text{dist}[u] + 1$ :  
         $\text{dist}[v] = \text{dist}[u] + 1$   
         $\text{pred}[v] = u$   
         $\text{Push}(Q, v)$ 
```

Recovering all paths from SSSP-BFS

A simple modification allows us to store all of the possible shortest paths.

We just need to adjust our $\text{pred}[u]$ data structure to store a *list* of predecessors, rather than just one!



SSSP-BFS(s):

$\text{dist}[u] \leftarrow \infty$ for all $u \in V$

$\text{pred}[u] \leftarrow \text{null}$ for all $u \in V$

$\text{dist}[s] \leftarrow 0$

$Q \leftarrow s$

While Q is not empty:

$u \leftarrow \text{Pull}(Q)$

For $v \in \text{Neighbors}(u)$:

If $\text{dist}[v] > \text{dist}[u] + 1$:

$\text{dist}[v] = \text{dist}[u] + 1$

$\text{pred}[v] = [u]$

$\text{Push}(Q, v)$

Else If $\text{dist}[v] = \text{dist}[u] + 1$:

Append u to $\text{pred}[v]$

$\text{Push}(Q, v)$

All Pairs Shortest Paths with BFS

We have an algorithm for computing shortest paths from a single source node to every other node

We need the shortest paths for *all pairs* of nodes (the *all pairs shortest paths* problem)

One option: Just run SSSP-BFS from every node!

```
APSP-BFS( $G = (V, E)$ ):  
  For  $v \in V$ :  
    SSSP-BFS( $v$ )
```

Running time

For each of n nodes, we run a full BFS. BFS runs in $O(n + m)$ time. Therefore we have $O(n(n + m))$, or $O(n^2 + nm)$.

All Pairs Shortest Paths with BFS

We have an algorithm for computing shortest paths from a single source node to every other node

We need the shortest paths for *all pairs* of nodes (the *all pairs shortest paths* problem)

One option: Just run SSSP-BFS from every node!

```
APSP-BFS-Paths( $G = (V, E)$ ):
```

```
  For  $s \in V$ :
```

```
    SSSP-BFS( $s$ ) // fills  $\text{pred}[u] \forall u$ 
```

```
  For  $t \in V$ :
```

```
    If  $s \neq t$  and  $s > t$ :
```

```
       $\text{paths}[s, t] \leftarrow \text{RecoverPaths}(s, t)$ 
```

Running time

For each of n nodes, we run a full BFS. BFS runs in $O(n + m)$ time. Therefore we have $O(n(n + m))$, or $O(n^2 + nm)$.

RecoverPaths is $O(n)$, meaning the doubly nested loop is $O(n^3)$, regardless of SSSP-BFS.

Betweenness Centrality

Now we can compute
betweenness centrality:

$$B(u) = \sum_{s \neq t \neq u} \frac{\sigma_{st}(u)}{\sigma_{st}}$$

Where σ_{st} is the number of shortest paths between nodes s and t and $\sigma_{st}(u)$ is the number of those shortest paths that include u .

```
Betweenness( $G$ ):  
  paths  $\leftarrow$  APSP-BFS-Paths( $G$ )  
  For  $u \in V$ :  
    For  $s \in V$ :  
      For  $t \in V$ :  
        if  $s \neq t$ :  
          denominator  $\leftarrow$  |paths[s,t]|  
          numerator  $\leftarrow$  |paths[s,t] that contain  $u$ |  
           $b[u] += \frac{\text{numerator}}{\text{denominator}}$   
  
  Return  $b$ 
```

Betweenness Centrality

Now we can compute
betweenness centrality:

$$B(u) = \sum_{s \neq t \neq u} \frac{\sigma_{st}(u)}{\sigma_{st}}$$

Where σ_{st} is the number of shortest paths between nodes s and t and $\sigma_{st}(u)$ is the number of those shortest paths that include u .

```
Betweenness( $G$ ):  
  paths  $\leftarrow$  APSP-BFS-Paths( $G$ )  
  For  $u \in V$ :  
    For  $s \in V$ :  
      For  $t \in V$ :  
        if  $s \neq t$ :  
          denominator  $\leftarrow$  |paths[s,t]|  
          numerator  $\leftarrow$  |paths[s,t] that contain  $u$ |  
           $b[u] += \frac{\text{numerator}}{\text{denominator}}$   
  
  Return  $b$ 
```

Question: Does all of this work on directed graphs?

Betweenness Centrality

Now we can compute
betweenness centrality:

$$B(u) = \sum_{s \neq t \neq u} \frac{\sigma_{st}(u)}{\sigma_{st}}$$

Where σ_{st} is the number of shortest paths between nodes s and t and $\sigma_{st}(u)$ is the number of those shortest paths that include u .

```
Betweenness( $G$ ):  
  paths  $\leftarrow$  APSP-BFS-Paths( $G$ )  
  For  $u \in V$ :  
    For  $s \in V$ :  
      For  $t \in V$ :  
        if  $s \neq t$ :  
          denominator  $\leftarrow$  |paths[ $s, t$ ] |  
          numerator  $\leftarrow$  |paths[ $s, t$ ] that contain  $u$  |  
           $b[u] += \frac{\text{numerator}}{\text{denominator}}$   
  
  Return  $b$ 
```

Question: Does all of this work on directed graphs?

Yes! With some modification to APSP-BFS-Paths to account for when a path does not exist ($\text{dist}[s,t]=\infty$)

What about weighted graphs?

So far, we have only considered unweighted graphs, or equivalently graphs with uniform weights.

We may want to find shortest paths in a *weighted graph* $G = (V, E, W)$ where W is a set of weights corresponding to the edges, e.g. $W = (u, v, w)$ where w is a nonnegative integer for all $(u, v) \in E$.

Generalizing SSSP-BFS: Best First Search

We can modify our BFS based algorithm to take edge weight into account

The distance corresponding to a path between two nodes is now the sum of the edge weights along the path

Modification requires taking a “global” view of the graph – the next step in any traversal algorithm involves choosing an edge to follow, we will choose it in a smarter way.

Dijkstra's Algorithm: choose the minimum distance edge to try to update next using a priority queue

Dijkstra's algorithm is an example of a “best first search” approach to graph traversal: we have some criteria (known as a heuristic) for choosing a good next node, so we use it.

SSSP-Dijkstra(s):

$\text{dist}[u] \leftarrow \infty$ for all $u \in V$

$\text{pred}[u] \leftarrow \text{null}$ for all $u \in V$

$\text{dist}[s] \leftarrow 0$

$Q \leftarrow (s, 0)$

While Q is not empty:

$u \leftarrow \text{PullMinimum}(Q)$

For $v \in \text{Neighbors}(u)$:

If $\text{dist}[v] > \text{dist}[u] + 1$:

$\text{dist}[v] = \text{dist}[u] + 1$

$\text{pred}[v] = [u]$


PushOrReplace($Q, v, \text{dist}[v]$)

Else If $\text{dist}[v] = \text{dist}[u] + 1$:

Append u to $\text{pred}[v]$

PushOrReplace($Q, v, \text{dist}[v]$)


Dijkstra's Algorithm: Demo



Dijkstra's Algorithm: Demo

Initialize

	A	B	C	D	E
$d_0(u)$	0	∞	∞	∞	∞




$$S = \{\}$$

Dijkstra's Algorithm: Demo

Explore A

	A	B	C	D	E
$d_0(u)$	0	∞	∞	∞	∞
$d_1(u)$	0	10	3	∞	∞




$$S = \{A\}$$

Dijkstra's Algorithm: Demo

Explore C


	A	B	C	D	E
$d_0(u)$	0	∞	∞	∞	∞
$d_1(u)$	0	10	3	∞	∞
$d_2(u)$	0	7	3	11	5



$$S = \{A, C\}$$

Dijkstra's Algorithm: Demo

Explore E




	A	B	C	D	E
$d_0(u)$	0	∞	∞	∞	∞
$d_1(u)$	0	10	3	∞	∞
$d_2(u)$	0	7	3	11	5
$d_3(u)$	0	7	3	11	5

$$S = \{A, C, E\}$$

Dijkstra's Algorithm: Demo


Explore B



	A	B	C	D	E
$d_0(u)$	0	∞	∞	∞	∞
$d_1(u)$	0	10	3	∞	∞
$d_2(u)$	0	7	3	11	5
$d_3(u)$	0	7	3	11	5
$d_4(u)$	0	7	3	9	5

Dijkstra's Algorithm: Demo

Don't need to explore D




	A	B	C	D	E
$d_0(u)$	0	∞	∞	∞	∞
$d_1(u)$	0	10	3	∞	∞
$d_2(u)$	0	7	3	11	5
$d_3(u)$	0	7	3	11	5
$d_4(u)$	0	7	3	9	5

$$S = \{A, C, E, B, D\}$$

Dijkstra's Algorithm: Demo

Maintain parent pointers so we can find the shortest paths



	A	B	C	D	E
$d_0(u)$	0	∞	∞	∞	∞
$d_1(u)$	0	10	3	∞	∞
$d_2(u)$	0	7	3	11	5
$d_3(u)$	0	7	3	11	5
$d_4(u)$	0	7	3	9	5

Dijkstra: Why does it work?

At the beginning, we have only that the distance from s to itself is 0, which is true by assumption.

SSSP-Dijkstra(s):

$\text{dist}[u] \leftarrow \infty$ for all $u \in V$

$\text{pred}[u] \leftarrow \text{null}$ for all $u \in V$

$\text{dist}[s] \leftarrow 0$

$Q \leftarrow (s, 0)$

While Q is not empty:

$u \leftarrow \text{PullMinimum}(Q)$

For $v \in \text{Neighbors}(u)$:

If $\text{dist}[v] > \text{dist}[u] + 1$:

$\text{dist}[v] = \text{dist}[u] + 1$

$\text{pred}[v] = [u]$

PushOrReplace($Q, v, \text{dist}[v]$)

Else If $\text{dist}[v] = \text{dist}[u] + 1$:

Append u to $\text{pred}[v]$

PushOrReplace($Q, v, \text{dist}[v]$)

Dijkstra: Why does it work?

At the beginning, we have only that the distance from s to itself is 0, which is true by assumption.

First, we explore the neighborhood of s and set all distances to its neighbors correctly.

SSSP-Dijkstra(s):

$\text{dist}[u] \leftarrow \infty$ for all $u \in V$

$\text{pred}[u] \leftarrow \text{null}$ for all $u \in V$

$\text{dist}[s] \leftarrow 0$

$Q \leftarrow (s, 0)$

While Q is not empty:

$u \leftarrow \text{PullMinimum}(Q)$

For $v \in \text{Neighbors}(u)$:

If $\text{dist}[v] > \text{dist}[u] + 1$:

$\text{dist}[v] = \text{dist}[u] + 1$

$\text{pred}[v] = [u]$

PushOrReplace($Q, v, \text{dist}[v]$)

Else If $\text{dist}[v] = \text{dist}[u] + 1$:

Append u to $\text{pred}[v]$

PushOrReplace($Q, v, \text{dist}[v]$)

Dijkstra: Why does it work?

At the beginning, we have only that the distance from s to itself is 0, which is true by assumption.

First, we explore the neighborhood of s and set all distances to its neighbors correctly.

Then we choose another node, call it v_1 , and correctly set all of the distances from $s \rightarrow v_1 \rightarrow t$, where $t \in \text{Neighbors}(v_1)$.

SSSP-Dijkstra(s):

$\text{dist}[u] \leftarrow \infty$ for all $u \in V$

$\text{pred}[u] \leftarrow \text{null}$ for all $u \in V$

$\text{dist}[s] \leftarrow 0$

$Q \leftarrow (s, 0)$

While Q is not empty:

$u \leftarrow \text{PullMinimum}(Q)$

For $v \in \text{Neighbors}(u)$:

If $\text{dist}[v] > \text{dist}[u] + 1$:

$\text{dist}[v] = \text{dist}[u] + 1$

$\text{pred}[v] = [u]$

PushOrReplace($Q, v, \text{dist}[v]$)

Else If $\text{dist}[v] = \text{dist}[u] + 1$:

Append u to $\text{pred}[v]$

PushOrReplace($Q, v, \text{dist}[v]$)

Dijkstra: Why does it work?

At the beginning, we have only that the distance from s to itself is 0, which is true by assumption.

First, we explore the neighborhood of s and set all distances to its neighbors correctly.

Then we choose another node, call it v_1 , and correctly set all of the distances from $s \rightarrow v_1 \rightarrow t$, where $t \in \text{Neighbors}(v_1)$.

At the i^{th} node, we set the correct distances from $s \rightsquigarrow v_i \rightsquigarrow t$, where $t \in \text{Neighbors}(v_i)$.

SSSP-Dijkstra(s):

$\text{dist}[u] \leftarrow \infty$ for all $u \in V$

$\text{pred}[u] \leftarrow \text{null}$ for all $u \in V$

$\text{dist}[s] \leftarrow 0$

$Q \leftarrow (s, 0)$

While Q is not empty:

$u \leftarrow \text{PullMinimum}(Q)$

For $v \in \text{Neighbors}(u)$:

If $\text{dist}[v] > \text{dist}[u] + 1$:

$\text{dist}[v] = \text{dist}[u] + 1$

$\text{pred}[v] = [u]$

PushOrReplace($Q, v, \text{dist}[v]$)

Else If $\text{dist}[v] = \text{dist}[u] + 1$:

Append u to $\text{pred}[v]$

PushOrReplace($Q, v, \text{dist}[v]$)

Dijkstra: Why does it work?

At the beginning, we have only that the distance from s to itself is 0, which is true by assumption.

First, we explore the neighborhood of s and set all distances to its neighbors correctly.

Then we choose another node, call it v_1 , and correctly set all of the distances from $s \rightarrow v_1 \rightarrow t$, where $t \in \text{Neighbors}(v_1)$.

At the i^{th} node, we set the correct distances from $s \rightsquigarrow v_i \rightsquigarrow t$, where $t \in \text{Neighbors}(v_i)$.

Invariant: After we explore the i^{th} node, $\text{dist}[u]$ is set correctly for all u visited so far

SSSP-Dijkstra(s):

$\text{dist}[u] \leftarrow \infty$ for all $u \in V$

$\text{pred}[u] \leftarrow \text{null}$ for all $u \in V$

$\text{dist}[s] \leftarrow 0$

$Q \leftarrow (s, 0)$

While Q is not empty:

$u \leftarrow \text{PullMinimum}(Q)$

For $v \in \text{Neighbors}(u)$:

 If $\text{dist}[v] > \text{dist}[u] + 1$:

$\text{dist}[v] = \text{dist}[u] + 1$

$\text{pred}[v] = [u]$


 PushOrReplace($Q, v, \text{dist}[v]$)

 Else If $\text{dist}[v] = \text{dist}[u] + 1$:

 Append u to $\text{pred}[v]$

 PushOrReplace($Q, v, \text{dist}[v]$)


Dijkstra: Why does it work?



Invariant: After we explore the i^{th} node, $\text{dist}[u]$ is set correctly for all u visited so far

We want to prove that $d_i(v) = d_i(u) + w_{uv}$ is the shortest path from s to v if v is the next node in the priority queue. We showed this works for $i = 1$ and $i = 2$.


Dijkstra: Why does it work?



Invariant: After we explore the i^{th} node, $\text{dist}[u]$ is set correctly for all u visited so far

We want to prove that $d_i(v) = d_i(u) + w_{uv}$ is the shortest path from s to v if v is the next node in the priority queue. We showed this works for $i = 1$ and $i = 2$.

Dijkstra: Why does it work?




Invariant: After we explore the i^{th} node, $\text{dist}[u]$ is set correctly for all u visited so far

We want to prove that $d_i(v) = d_i(u) + w_{uv}$ is the shortest path from s to v if v is the next node in the priority queue. We showed this works for $i = 1$ and $i = 2$.

Consider the picture above, which represents two possibilities for paths from s to some node v . The path P represents an actual shortest path, while P' represents an alternative path assuming some node y was actually a better next choice than v , meaning that $\ell(P') < \ell(P)$. We have:

Dijkstra: Why does it work?




Invariant: After we explore the i^{th} node, $\text{dist}[u]$ is set correctly for all u visited so far

We want to prove that $d_i(v) = d_i(u) + w_{uv}$ is the shortest path from s to v if v is the next node in the priority queue. We showed this works for $i = 1$ and $i = 2$.

Consider the picture above, which represents two possibilities for paths from s to some node v . The path P represents an actual shortest path, while P' represents an alternative path assuming some node y was actually a better next choice than v , meaning that $\ell(P') < \ell(P)$. We have:

$$\ell(P') = d_i(x) + w_{xy} + w_{yv}$$

Dijkstra: Why does it work?




Invariant: After we explore the i^{th} node, $\text{dist}[u]$ is set correctly for all u visited so far

We want to prove that $d_i(v) = d_i(u) + w_{uv}$ is the shortest path from s to v if v is the next node in the priority queue. We showed this works for $i = 1$ and $i = 2$.

Consider the picture above, which represents two possibilities for paths from s to some node v . The path P represents an actual shortest path, while P' represents an alternative path assuming some node y was actually a better next choice than v , meaning that $\ell(P') < \ell(P)$. We have:

$$\begin{aligned}\ell(P') &= d_i(x) + w_{xy} + w_{yv} \\ &\geq d_i(x) + w_{xy} \quad \text{Since } w_{yv} \geq 0\end{aligned}$$

Dijkstra: Why does it work?



Invariant: After we explore the i^{th} node, $\text{dist}[u]$ is set correctly for all u visited so far

We want to prove that $d_i(v) = d_i(u) + w_{uv}$ is the shortest path from s to v if v is the next node in the priority queue. We showed this works for $i = 1$ and $i = 2$.


Consider the picture above, which represents two possibilities for paths from s to some node v . The path P represents an actual shortest path, while P' represents an alternative path assuming some node y was actually a better next choice than v , meaning that $\ell(P') < \ell(P)$. We have:

$$\begin{aligned}\ell(P') &= d_i(x) + w_{xy} + w_{yv} \\ &\geq d_i(x) + w_{xy} \\ &\geq d_i(y)\end{aligned}$$

Since $w_{yv} \geq 0$

We know x is explored already

Dijkstra: Why does it work?



Invariant: After we explore the i^{th} node, $\text{dist}[u]$ is set correctly for all u visited so far

We want to prove that $d_i(v) = d_i(u) + w_{uv}$ is the shortest path from s to v if v is the next node in the priority queue. We showed this works for $i = 1$ and $i = 2$.

Consider the picture above, which represents two possibilities for paths from s to some node v . The path P represents an actual shortest path, while P' represents an alternative path assuming some node y was actually a better next choice than v , meaning that $\ell(P') < \ell(P)$. We have:


$$\begin{aligned}\ell(P') &= d_i(x) + w_{xy} + w_{yv} \\ &\geq d_i(x) + w_{xy} \\ &\geq d_i(y) \\ &\geq d_i(v)\end{aligned}$$

Since $w_{yv} \geq 0$

We know x is explored already

We chose v to explore, not y !

Dijkstra: Why does it work?



Invariant: After we explore the i^{th} node, $\text{dist}[u]$ is set correctly for all u visited so far

We want to prove that $d_i(v) = d_i(u) + w_{uv}$ is the shortest path from s to v if v is the next node in the priority queue. We showed this works for $i = 1$ and $i = 2$.

Consider the picture above, which represents two possibilities for paths from s to some node v . The path P represents an actual shortest path, while P' represents an alternative path assuming some node y was actually a better next choice than v , meaning that $\ell(P') < \ell(P)$. We have:

$$\begin{aligned}\ell(P') &= d_i(x) + w_{xy} + w_{yv} \\ &\geq d_i(x) + w_{xy} \\ &\geq d_i(y) \\ &\geq d_i(v) \\ &= \ell(P)\end{aligned}$$

Since $w_{yv} \geq 0$

We know x is explored already

We chose v to explore, not y !

So $\ell(P') \geq \ell(P)$. Contradiction!

Dijkstra running time

Assuming our priority queue supports insertion, update, and extraction in $O(\log E)$ time, this approach runs in $O(n + \log E)$

SSSP-Dijkstra(s):

$\text{dist}[u] \leftarrow \infty$ for all $u \in V$

$\text{pred}[u] \leftarrow \text{null}$ for all $u \in V$

$\text{dist}[s] \leftarrow 0$

$Q \leftarrow (s, 0)$

While Q is not empty:

$u \leftarrow \text{PullMinimum}(Q)$

For $v \in \text{Neighbors}(u)$:

If $\text{dist}[v] > \text{dist}[u] + 1$:

$\text{dist}[v] = \text{dist}[u] + 1$

$\text{pred}[v] = [u]$

PushOrReplace($Q, v, \text{dist}[v]$)

Else If $\text{dist}[v] = \text{dist}[u] + 1$:

Append u to $\text{pred}[v]$

PushOrReplace($Q, v, \text{dist}[v]$)

Floyd-Warshall

What about applications where negative edgeweights make sense?

- Transactions
- Chemical reactions
- Changes over time

The *Floyd-Warshall* algorithm is a dynamic programming solution to solving the all-pairs-shortest-paths problem on weighted, directed graphs that have no *negative cycles*.

Next Time

Spanning trees and flow algorithms

Suggested Reading: Erickson Chapter 7 and Chapter 10 through 10.3